

# Processing In Memory

## Dual Degree Stage 1 Report

Submitted in partial fulfillment of the requirements  
for the

Dual Degree Programme

by

**Nayan Barhate**  
(Roll No. 180070037)

Under the guidance of  
**Prof. Virendra Singh**



Department of Electrical Engineering  
Indian Institute of Technology Bombay  
October 2022

## **Acknowledgement**

I express my gratitude to my guide Prof. Virendra Singh for providing me the opportunity to work on this topic.

Nayan Barhate  
Electrical Engineering  
IIT Bombay

## **Abstract**

Processing-in-memory (PIM) is rapidly rising as a viable solution for the memory wall crisis, rebounding from its unsuccessful attempts in 1990s due to practicality concerns, which are alleviated with recent advances in 3D stacking technologies. However, it is still challenging to integrate the PIM architectures with existing systems in a seamless manner due to two common characteristics: unconventional programming models for in-memory computation units and lack of ability to utilize large on-chip caches.

# Contents

<b>List of Figures</b>	<b>2</b>
<b>1 Introduction</b>	<b>4</b>
<b>2 Literature Survey</b>	<b>5</b>
2.1 PIM Enabled Instruction . . . . .	5
2.1.1 PIM Compute Unit(PCU) . . . . .	5
2.1.2 PIM Management Unit(PMU) . . . . .	6
2.1.3 PEI execution . . . . .	6
2.1.4 Results . . . . .	8
2.1.5 Conclusion . . . . .	9
<b>3 Future Work</b>	<b>10</b>

# List of Figures

1.1	Performance improvement with an in-memory atomic addition operation used for the PageRank algorithm. . . . .	4
2.1	Overview of the proposed architecture. . . . .	5
2.2	Host-side PEI execution . . . . .	7
2.3	Memory-side PEI execution. . . . .	7
2.4	Speedup comparison under different input sizes. . . . .	8

# Chapter 1

## Introduction

All computing systems, including cloud and server platforms, desktop computers, mobile and embedded devices, and sensors, rely heavily on main memory. It is one of the primary pillars of any computing platform, along with 1) processing elements (or computational elements), which can include CPU cores, GPU cores, accelerators, or re-configurable devices, and 2) communication elements, which can include interconnects, network interfaces, and network processing units.

A simple addition operation when replaced by an atomic addition command inside memory gives significant speedup(53%) to graphs with high number of vertices(5M), but the same command gives us a speeddown(upto 20%) when executed on graphs with low number of vertices(62K). This performance degradation is attributed to the observation that sometimes the data on which the command was executed is readily available in cache. Thus PIM is not always the panacea and the decision of 'to PIM or Not' needs to be handled based on the data locality of the operands.




Figure 1.1: Performance improvement with an in-memory atomic addition operation used for the PageRank algorithm.

# Chapter 2

## Literature Survey

### 2.1 PIM Enabled Instruction

For easier adoption of PIM, we desire little to no change in programmer's model. Main memory products with integrated full-fledged processors with new programming models may also not be available in the near future because of the associated design complexity and changes required across the hardware/software stack. Also, prior proposals do not utilize the benefits of on-chip caches and virtual memory. This paper overcomes these challenges by extending the ISA of the host processor with PIM-Enabled Instructions (PEIs)[1], without changing the existing programming model.

The paper proposes simple ISA extensions and the architectural support required to integrate such simple operations into conventional systems. The architecture proposes a restriction on the memory region accessible by a single PIM operation is limited to a single last-level cache block. This ensures the data can be accessed by the same vertical link and simplifies the locality profiling.




Figure 2.1: Overview of the proposed architecture.

#### 2.1.1 PIM Compute Unit(PCU)

PCU is a hardware unit that executes PEIs. All PCUs in the system have the same computation logic so that any PEI can be executed by any PCU. In the proposed architecture one PCU is placed besides every core in the host processor and one PCU is placed in every vault. PCU has a SRAM based operand buffer to store information of in-flight PEIs.

## 2.1.2 PIM Management Unit(PMU)

The function of PMU is to

- Ensure atomicity among different PEIs
- Maintain cache coherency between memory and cache during PIM operation
- Profile data locality for locality aware execution

### PIM Directory

PIM Directory ensures atomicity among different in-flight PEIs. Atomicity between a PEI and normal memory instruction is not maintained by PIM Directory. This must be maintained by the programmer by using pfence. PIM Directory is a directly mapped, tag-less table indexed by XOR-folded addresses of target cache blocks. Each entry implements a reader-writer pair lock

### Cache Coherency Management

Since the PEIs are resisted to a single last-level cache block, the cache coherency problem is simplified. Whenever the PEI is executed on the host processor the corresponding last level block is flushed out and subsequent blocks in L1 and L2 are back-invalidated. Thus ensuring neither the memory nor the processor has a stale copy of data.

### Locality Monitor

The key idea behind locality monitor is to decide whether to execute the PEI locally on the host or remotely in the memory. Locality monitor is a tag array with same number of sets/ways as that of the last-level cache. The key difference between the locality monitor and the tag array for the last-level cache is that the former is also updated when a PIM operation is issued to memory. In our locality monitor, the data locality of a PEI can be identified by checking to see if its target cache block address hits in the locality monitor.

## 2.1.3 PEI execution

### Host-side PEI execution

1. The out-of-order core writes the input operands to the memory-mapped registers of the PCU and issues the PEI
2. PIM Directory then checks for atomicity and data locality - which results in high locality of data indicating host-side execution
3. The PCU loads the block from cache and executes the instruction
4. After execution the PCU writes back the output operands in the cache and notifies the PMU for completion
5. The out-of-order core then reads the output operands using the memory-mapped registers of PCU.



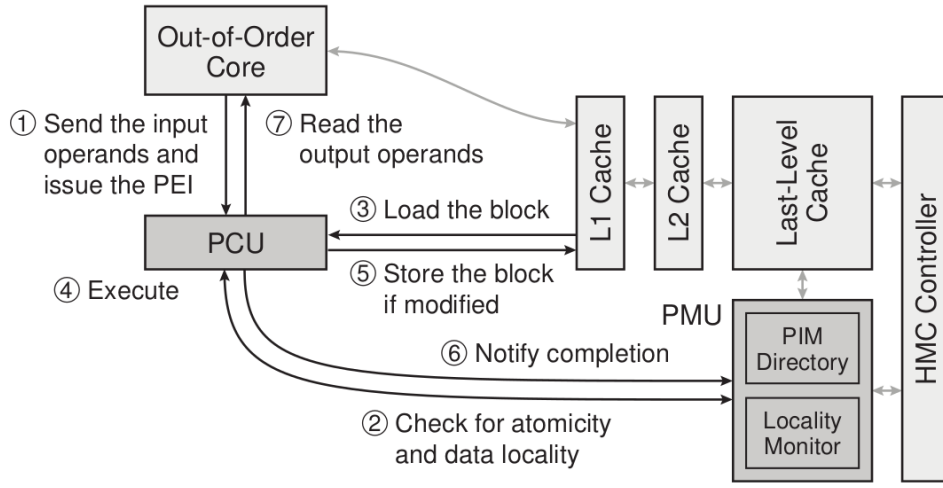


Figure 2.2: Host-side PEI execution

### Memory-side PEI execution

1. The out-of-order core writes the input operands to the memory-mapped registers of the PCU and issues the PEI
2. PIM Directory then checks for atomicity and data locality - which results in low locality of data indicating memory-side execution
3. The Last-level block in cache is flushed and the corresponding blocks in L1 and L2 are back-invalidated
4. The host side PCU sends the input operands to the memory side PCU using the HMC controller
5. The PCU executes the operation and sends back the output operands indicating completion
6. The out-of-order core then reads the output operands using the memory-mapped registers of PCU.

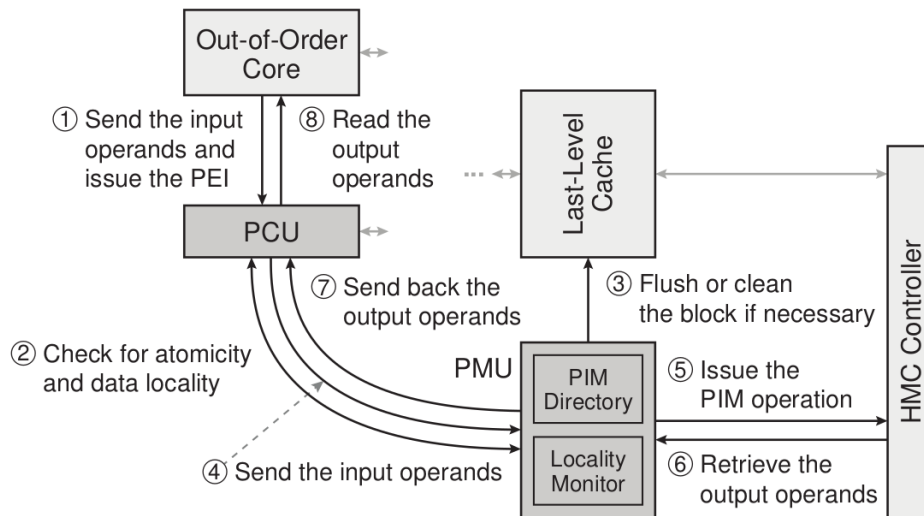


Figure 2.3: Memory-side PEI execution.

## 2.1.4 Results

- For large inputs, PIM-Only achieves 44% speedup over Ideal-Host since using PIM operations better utilizes the vertical DRAM bandwidth inside HMCs
- For small inputs, PIM-Only degrades average performance by 20% because PIM operations always access DRAM even though the data set comfortably fits in on-chip caches
- Locality-Aware system provides the speedup of PIM-Only in workloads with large inputs (47% improvement over Host-Only) by offloading 79% of PEIs to memory-side PCUs.
- Locality-Aware system provides the speedup of PIM-Only in workloads with small inputs (32% improvement over Host-Only) by offloading 86% of PEIs to memory-side PCUs.
- Locality-Aware often outperforms both Host-Only and PIM-Only by simultaneously utilizing host-side and memory-side PCUs for PEI execution




Figure 2.4: Speedup comparison under different input sizes.

### 2.1.5 Conclusion

- PIM-enabled instructions (PEIs) is a practical model for processing-in-memory and its hardware implementation, which is compatible with existing cache coherence and virtual memory mechanisms
- PEIS have a last-level cache block restriction to ensure atomicity.
- PEIs use data locality to decide the processing core.

# Chapter 3

## Future Work

I plan to read the following conference papers -

- A. Devic et al., "To PIM or not for emerging general purpose processing in DDR memory systems". In Proceedings of the 49th Annual International Symposium on Computer Architecture (ISCA '22). Association for Computing Machinery, New York, NY, USA, 231–244.
- K. Hsieh et al., "Transparent Offloading and Mapping (TOM): Enabling Programmer-Transparent Near-Data Processing in GPU Systems," 2016 ACM/IEEE 43rd Annual International Symposium on Computer Architecture (ISCA), 2016, pp. 204-216,

# References

- [1] J. Ahn, S. Yoo, O. Mutlu, and K. Choi, “Pim-enabled instructions: A low-overhead, locality-aware processing-in-memory architecture,” *SIGARCH Comput. Archit. News*, vol. 43, p. 336–348, jun 2015.